

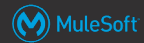


Module 7

Architecting and Deploying Effective API Implementations

1

At the end of this module, you should be able to



- Describe **auto-discovery** of API implementations implemented as Mule apps
- Appreciate how **Connectors** serve System APIs
- Describe **CloudHub**'s features and technology architecture
- Choose **Object Store** in a CloudHub setting
- Apply strategies that help API clients guard against **failures in API invocations**
- Describe the role of **CQRS** and the separation of commands and queries in API-led connectivity
- Explain the role of **Event Sourcing**

All contents © MuleSoft Inc.

2

2

Section 1

Developing API implementations for Mule runtimes and deploying them to CloudHub



3

Introducing API implementations management on Anypoint Platform



API implementations may be

- Developed as **Mule apps** for the Mule runtime
 - **Studio** or **Flow designer**
 - Deployed to an Anypoint Platform **runtime plane** such as the **CloudHub**
 - API Management, Analytics, etc. provided by an Anypoint Platform control plane **with or without separate API proxies**
- Developed **for other runtimes** (Node.js, Spring Boot, etc.)
 - Deployed to matching runtimes outside Anypoint Platform
 - API Management, Analytics, etc. provided by an Anypoint Platform control plane via **API proxies executing in a runtime plane**

All contents © MuleSoft Inc.

4

4

Introducing Auto-discovery of API implementations



- Auto-discovery is
 - **Auto-registration** of an API implementation
 - Developed as Mule app and deployed to Mule runtime (incl. API proxies)
 - With **API Manager**
- **On start-up** registers with API Manager as an implementation of a given API instance (implying API, version and environment)
- **Receives all API policies** for that API instance
 - By default refuses API invocations until all API policies have been applied
 - Gatekeeper feature

All contents © MuleSoft Inc.

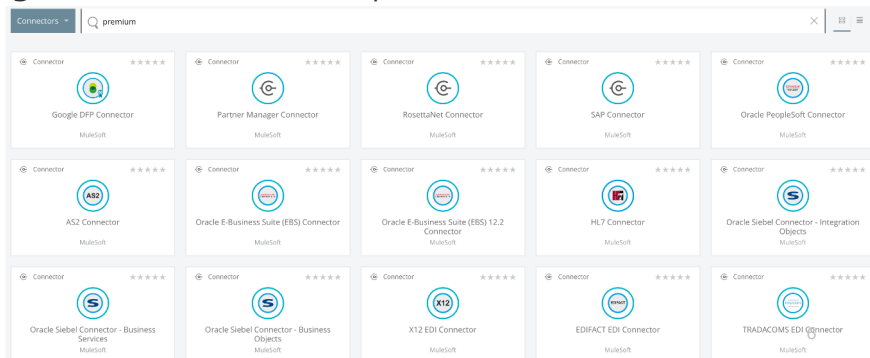
5

5

Anypoint Connectors to implement System APIs



- Connectors are **pre-built** (Anypoint Connectors) or **custom-developed** components that integrate Mule apps with external systems
 - More than **120** pre-built, many bundled with **Studio**
 - Others in **Exchange** and MuleSoft **Nexus** repo
 - **Support** categories:
 - Premium
 - Select
 - Certified
 - Community




All contents © MuleSoft Inc.

6

Connectors to implement System APIs




- Mainframe-based **Policy Admin System**
 - Accessed via **WebSphere MQ** using **JMS Connector**
- WebSphere-deployed **Motor Claims System**
 - Directly accessing its **DB/2** database using **Database Connector**
- Web-based **Home Claims System**
 - Exposes **SOAP APIs** accessed using **Web Service Consumer**



7

7

Fundamentals of CloudHub Technology Architecture



8

8

CloudHub worker sizes



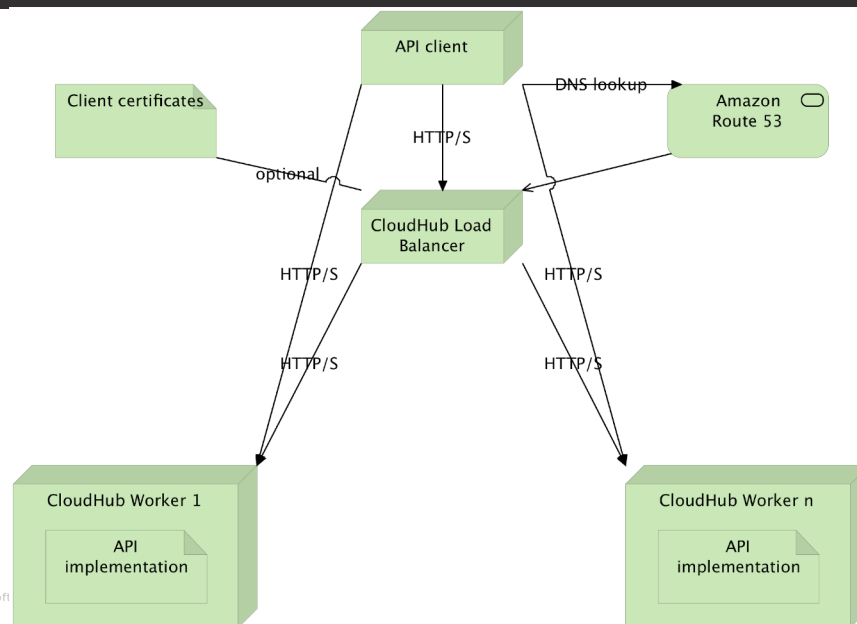
Worker Name	Heap Memory	Storage
0.1 vCores	500 MB	8 GB
0.2 vCores	1 GB	8 GB
1 vCores	1.5 GB	12 GB
2 vCores	3.5 GB	40 GB
4 vCores	7.5 GB	88 GB
8 vCores	15 GB	168 GB
16 vCores	32 GB	328 GB

All contents © MuleSoft Inc.

9

9

Fundamentals of CloudHub Technology Architecture




All contents © MuleSoft

10

10

Performance of CloudHub workers of various sizes



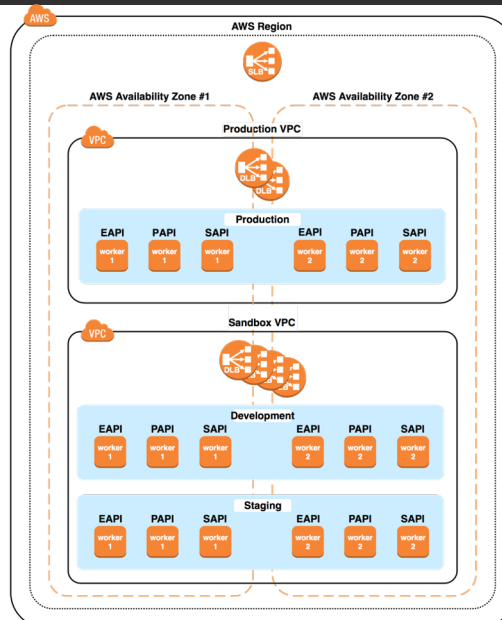
- Throughput as HTTP requs/s
- of a simple API implementation

All contents © MuleSoft Inc.

11

11

CloudHub Shared Worker Cloud and Anypoint VPC



All contents © MuleSoft Inc.

12

12

CloudHub Shared Worker Cloud and Anypoint VPC

EnvironmentsBusiness GroupsFirewall RulesInternal DNS

Learn more about [Firewall Rules](#)

[⊕ Add New Rule](#)

Type		Source		Port Range	
https.private.port	▼	Local VPC (10.2.0.0/16)	▼	8092	...
https.port	▼	Anywhere (0.0.0.0/0)	▼	8082	...
http.private.port	▼	Local VPC (10.2.0.0/16)	▼	8091	...
http.port	▼	Anywhere (0.0.0.0/0)	▼	8081	...

All contents © MuleSoft Inc.

13

13

Anypoint VPC connectivity: on-premises networks


The diagram illustrates the connectivity between a CloudHub VPC and an on-premises data center. On the left, the CloudHub VPC [10.0.0.0/24] contains four worker nodes. On the right, the On-Prem Data Center [192.168.177.0/24] contains two database instances. A dashed line labeled 'IPsec Tunnel' connects the two networks. Below the tunnel, the CloudHub External Gateway IP Address is [54.208.209.9] and the Your External Gateway IP Address is [217.18.13.66].

All contents © MuleSoft Inc.

14

14

Anypoint VPC connectivity: AWS VPC peering

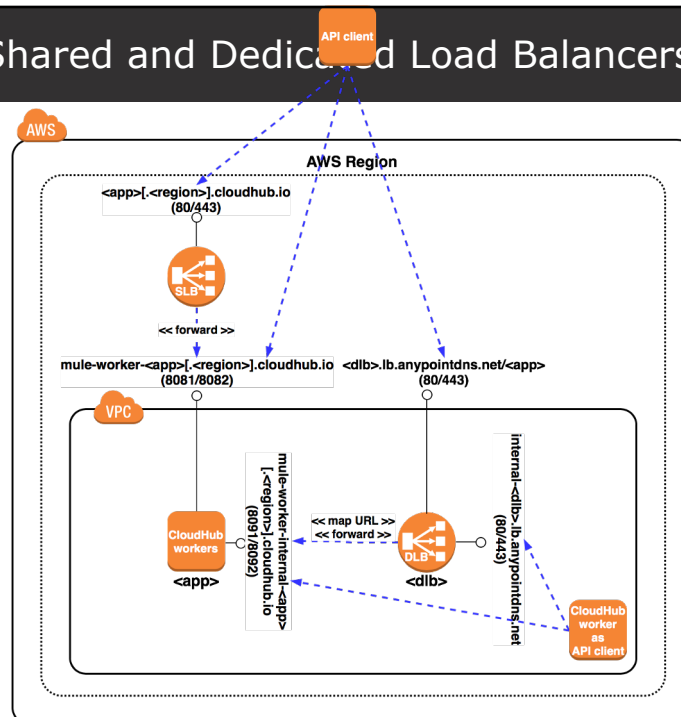


All contents © MuleSoft Inc.

15

15

CloudHub Shared and Dedicated Load Balancers



All contents © MuleSoft Inc.

16

16

Certificates for SLB, DLB and Mule applications



- Certificates will be required for **HTTPS** traffic between
 - API clients to CloudHub load balancer (SLB, DLB)
 - CloudHub load balancer to CH workers (Mule apps, API implemenations)
- **SLB certificates** are managed by **MuleSoft** and cannot be configured by customers
- **DLB certificate** management is the **customer's** responsibility
 - Use a certificate issued by a recognized **certificate authority**
 - Do not use a self-signed certificate
- CH-deployed **Mule apps** should use a **self-signed certificate**
 - CH workers' DNS names in **cloudhub.io** domain owned by MuleSoft

All contents © MuleSoft Inc.

17

17

CloudHub DNS names



	Control Plane Region	
	US East	EU (Frankfurt)
Web UI		
Platform APIs	anypoint.mulesoft.com	eu1.anypoint.mulesoft.com
DLB Public IPs	<dlb>.lb.anypointdns.net	<dlb>.lb-prod-eu-rt.anypointdns.net
DLB Private IPs	internal-<dlb>.lb.anypointdns.net	internal-<dlb>.lb-prod-eu-rt.anypointdns.net

All contents © MuleSoft Inc.

18

18

CloudHub DNS names for some runtime plane regions

mule-worker- mule- worker-internal-	Control Plane Region	
	US East	EU (Frankfurt)
US East	<app>.us-e1.cloudhub.io	N/A
EU (Frankfurt)	<app>.de-c1.cloudhub.io	<app>.de-c1.eu1.cloudhub.io
EU (Ireland)	<app>.ir-e1.cloudhub.io	<app>.ir-e1.eu1.cloudhub.io
US West (Oregon)	<app>.us-w2.cloudhub.io	N/A
EU (London)	<app>.uk-e1.cloudhub.io	N/A
Canada	<app>.ca-c1.cloudhub.io	N/A
Singapore	<app>.sg-s1.cloudhub.io	N/A
Sydney	<app>.au-s1.cloudhub.io	N/A

19

Exercise: CloudHub for Acme Insurance

Acme Insurance will deploy all their API implementations to **CloudHub**, under the **US control plane**. Based on your knowledge of their requirements (making assumptions if necessary) sketch their **Production environment**:

1. Region(s) of their Anypoint Platform runtime plane(s)
2. Number of workers per API implementation and worker size
3. Need for and use of Shared Worker Cloud, Anypoint VPCs, Shared/Dedicated Load Balancers, IPsec tunnels, VPC peering
4. API URL patterns, in particular hostnames, as seen by API clients
5. Ports on which API implementations listen

All contents © MuleSoft Inc.

20

20

Exercise: CloudHub for Acme Insurance

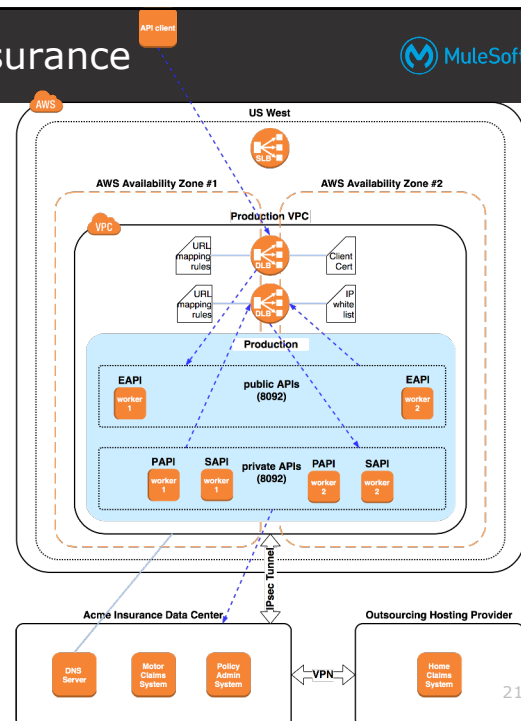


External API clients:

<https://api.acmeinsurance.com/aggregatortecreation/> ->
ans-aggregatortecreation-eapi

VPC-internal API clients:

<https://private-api.acmeinsurance.com/policyoptionsranking-papi/> ->
ans-policyoptionsranking-papi



All contents © MuleSoft Inc.

21

Understanding Object Store



- OS provides **key-value persistence** to all Mule apps
- Use for:
 - **correlation information** for async processing, circuit breakers, **caches**, ...
 - **idempotent** filters/validators, **watermarks**, **token stores**, stateful API policies such as **HTTP Caching**, ...
- Object Store **Connector**
- Mule app can use **several OS instances**
 - **No inter-Mule app data exchange**
- Values can also be **JSON** documents
- **No transactions or queries**
- **Persistent** or **transient**, **TTL**, max capacity
- **Default persistent OS** always available

All contents © MuleSoft Inc.

22

22

Understanding Object Store in CloudHub



- Keeps data in **same AWS region** as the owning Mule app
 - Available (only) to **all workers** of the owning Mule app
- OS content exposed in **Runtime Manager**
- Max **TTL of 30 days**
- **Unlimited** number of **keys**; each **value** limited to **10 MB**
- **Cross-AZ HA** within an Amazon Web Services region
- **Encryption** in transit and at rest

All contents © MuleSoft Inc.

23

23

Section 2

Designing API clients to use fault-tolerant API invocations



24

Exercise: Failures in API invocations



1. List every kind of failure that can occur when an API client invokes an API implementation
2. Think of ways of guarding against each of the failures you have identified

All contents © MuleSoft Inc.

25

25

Exercise: Failures in API invocations



The invocation of an API implementation by an API client fails if any of the **intervening components** fails at the moment of the API invocation:

- **Hardware**
- **Virtualization** or **container** stacks
- **Operating system, JVM** or other runtime and **libraries**
- Any **network component** such as Ethernet cards, switches, routers, cables, WIFI transmitters, ...
- Includes all **Anypoint Platform components** and underlying **AWS** services: API policies, API proxies, load balancers, CloudHub workers, etc.
- The **API implementation** itself, because of a bug in the app code, a failed deployment, ...

All contents © MuleSoft Inc.

26

26

Understanding the significance of failure in application networks



- High degree of **dependency** of any API implementation on other APIs
 - Inherent and **desired** feature of application networks
 - Shows high degree of **reuse** of APIs
- **Failures** in invocation of an API **affect many other APIs**
 - If unchecked, failures in API invocations propagate **transitively** through an application network
- Therefore: Make **API invocations fault-tolerant**

All contents © MuleSoft Inc.

27

27


Designing API clients for fault-tolerant API invocations



- **Fault-tolerant API invocation:**
 - **Invocation** from an API client to an API implementation via its API
 - Where **failure** of the API invocation
 - Does not necessarily lead to **failure** of the API client
 - API client is typically itself an API implementation
 - **Breaks transitive failure propagation**

- **Strategies:**

- Timeout
- Retry
- Circuit Breaker
- Fallback Invocation
- Opportunistic Parallel Invocations
- Fallback Result
- Client-Side Cache

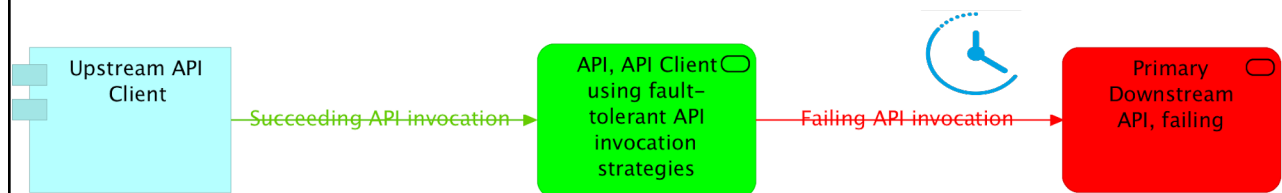


All contents © MuleSoft Inc.

28

28

Using timeouts to guard against slow APIs



All contents © MuleSoft Inc.


29

29

Using timeouts to guard against slow APIs



- **Slow API invocations**
 - Jeopardize **SLA** of API client at risk
 - Have higher probability of **eventually failing**
- Set **timeouts** for API invocations carefully
 - In accordance with the SLA of the API client
 - As **low** as possible, given SLA of invoked API
- Timed-out API invocation is indistinguishable from a **failed** API invocation
 - Requires other **additional** fault-tolerance strategies
- For **Mule apps**:
 - At the level of the **HTTP request**
 - For an entire **transaction**
 - As part of higher-level control constructs like **Scatter-Gather**



All contents © MuleSoft Inc.

30

30

Using timeouts to guard against slow APIs



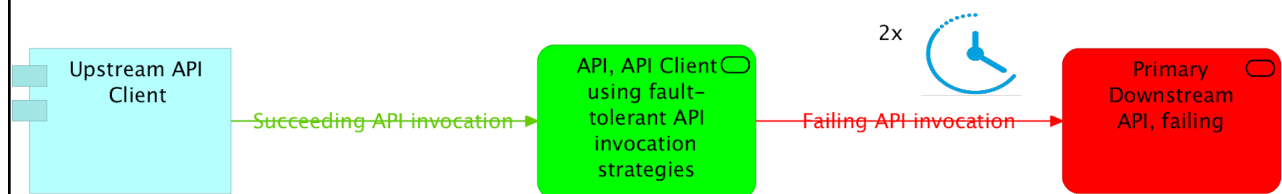
- "Aggregator Quote Creation EAPI" has SLA of median/max response time of **200 ms/500 ms**
- Synchronously invokes **3 Process APIs**, in turn, starting with "Policy Holder Search PAPI"
 - Task performed by "Policy Holder Search PAPI" is the simplest and **fastest** of all 3 Process APIs
 - "Aggregator Quote Creation EAPI" should define timeout of no more than approx. **100 ms** for the invocation of the "Policy Holder Search PAPI"
 - Most invocations (98%) must be guaranteed to complete within 100 ms

All contents © MuleSoft Inc.

31

31

Retrying failed API invocations



All contents © MuleSoft Inc.

32

32

Retrying failed API invocations



- Failed API invocation may **succeed if retried**
 - Only if failure was **transient**
- Difficult for API client to **identify** transient failures
- For **REST** APIs:
 - **4xx** response codes signify permanent failures, except
 - HTTP 408 Request Timeout
 - HTTP 423 Locked
 - HTTP 429 Too Many Requests
 - **5xx** response codes may signify transient failures
 - Only **idempotent** HTTP methods may be retried:
 - GET, HEAD, OPTIONS, PUT, DELETE

All contents © MuleSoft Inc.

33

33

Retrying failed API invocations



- To **limit** processing time:
 - Only **few** retries
 - **Short** wait times between retries
 - Short **timeouts** for each API invocation
- For **Mule apps**:
 - **HTTP Request Connector**
 - Has configurable support for interpreting HTTP response codes as failures
 - **Until Successful Scope**

All contents © MuleSoft Inc.

34

34

Invoking failing APIs less often with Circuit Breakers

- Failing API implementation may need time to **recover**
 - Continued invocations may **hinder** recovery
 - API client is **saved wasted time** and effort of invoking a failing API
- **Circuit Breaker**
 - **Monitors** API invocations
 - After "many" failed API invocations "**opens**"
 - Immediately fails invocations **without invoking** API
 - After "recovery period" **resumes** API invocations
 - Immediately "**opening**" upon failure
 - After "many" successful API invocations "**closes**"

All contents © MuleSoft Inc.

35

35

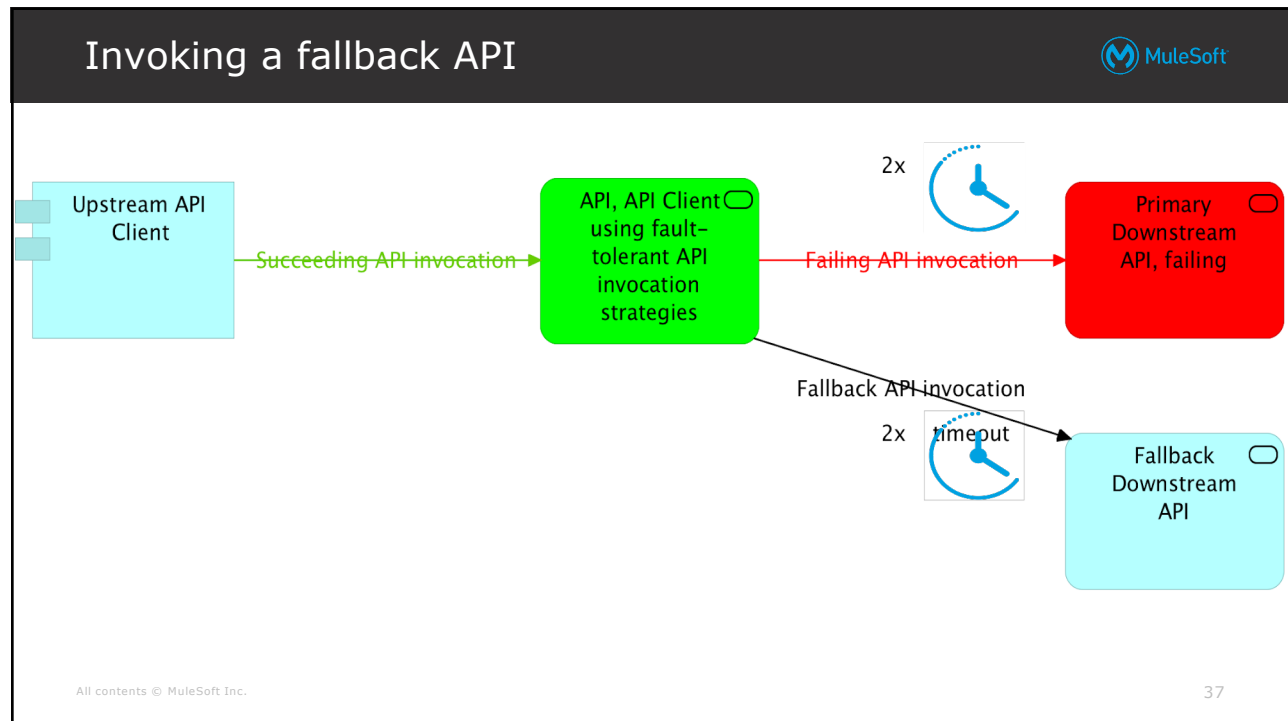
Invoking failing APIs less often with Circuit Breakers

- Circuit Breaker is a **stateful** component
 - Keeps track of health state of API **for API clients**
- Varying **scope**/reach of API clients that experience same state:
 - One API client in a **single CloudHub worker**
 - All instances of an API client app in **all CloudHub workers**
 - Requires remote communication between instances of the API client
 - All instances of **any API client** app in any CloudHub worker
 - Circuit Breaker as an application network-wide shared resource

All contents © MuleSoft Inc.

36

36



37

Invoking a fallback API

- After repeated API invocation failures it may be possible to invoke a **different API** as a fallback
 - May not be ideal but **sufficient**
 - Prefer **service degradation** over failure
- E.g., after repeated failures invoking "Motor Policy Holder Search SAPI" try:
 - Old but sufficiently compatible **version** of the same API
 - **Endpoint** of the same API from the DR site or other ClouHub region
 - "Motor Policy Search SAPI" which provides **functional super-set**
- For **Mule apps**:
 - **Until Successful Scope** plus exception strategies allow configuring fallback actions such as fallback API invocations

All contents © MuleSoft Inc.

38

Opportunistically invoking APIs in parallel



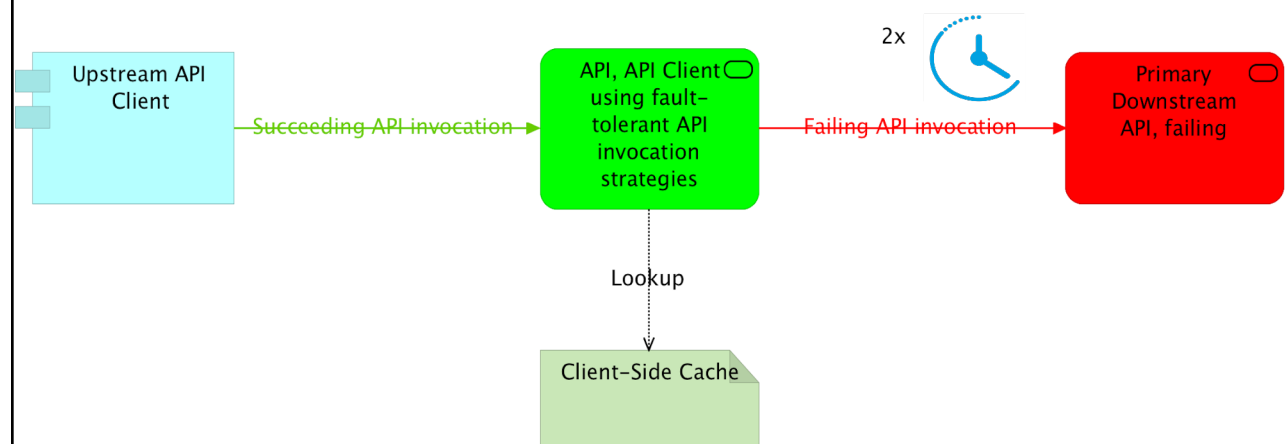
- Invocation of primary API and fallback API may be performed in **parallel**
 - **First response** to arrive is used, the other discarded
- Reduces overall **execution time** compared to serial fallback invocation
- Opportunistic, **egotistical** strategy
 - Puts increased **load** on the application network
 - Only in **exceptional** cases

All contents © MuleSoft Inc.

39

39

Using a previously cached result as a fallback



All contents © MuleSoft Inc.

40

40

Using a previously cached result as a fallback



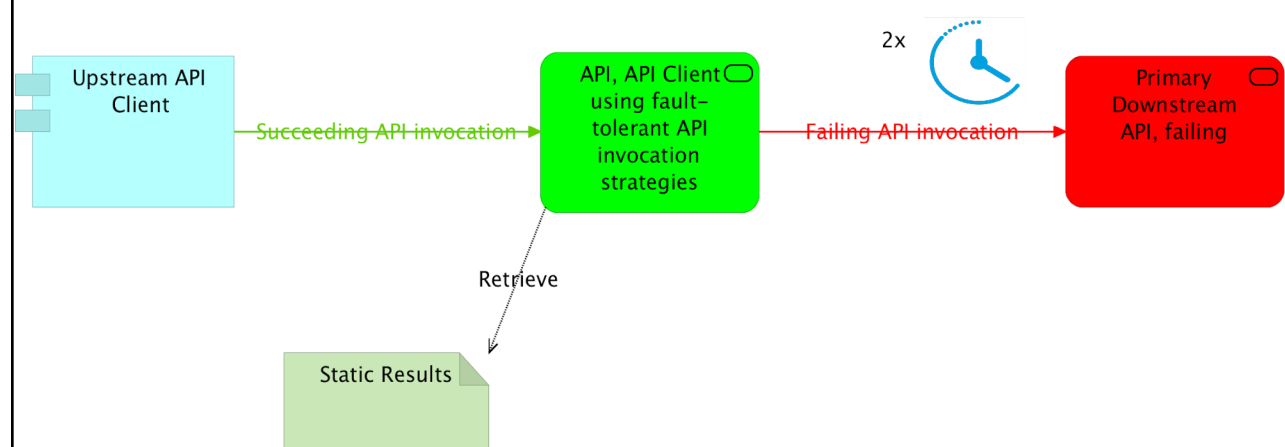
- **Client-side caching** by the API client is a great source of fallback results
- Governed by the usual **HTTP caching** rules:
 - Only responses from **safe** HTTP methods should be cached:
 - GET, HEAD, OPTIONS
 - HTTP **caching headers** should be honored (but may not)
 - Cache-Control, Last-Modified, Age, ...
- Increases **memory** footprint and processing of API client
- E.g., cached response from previous "Policy Options Retrieval SAPI" invocations may be used in case of failure
- For **Mule apps**:
 - **Cache Scope** and **Object Store Connector** support client-side caching

All contents © MuleSoft Inc.

41

41

Using a static fallback result



All contents © MuleSoft Inc.

42

42

Using a static fallback result



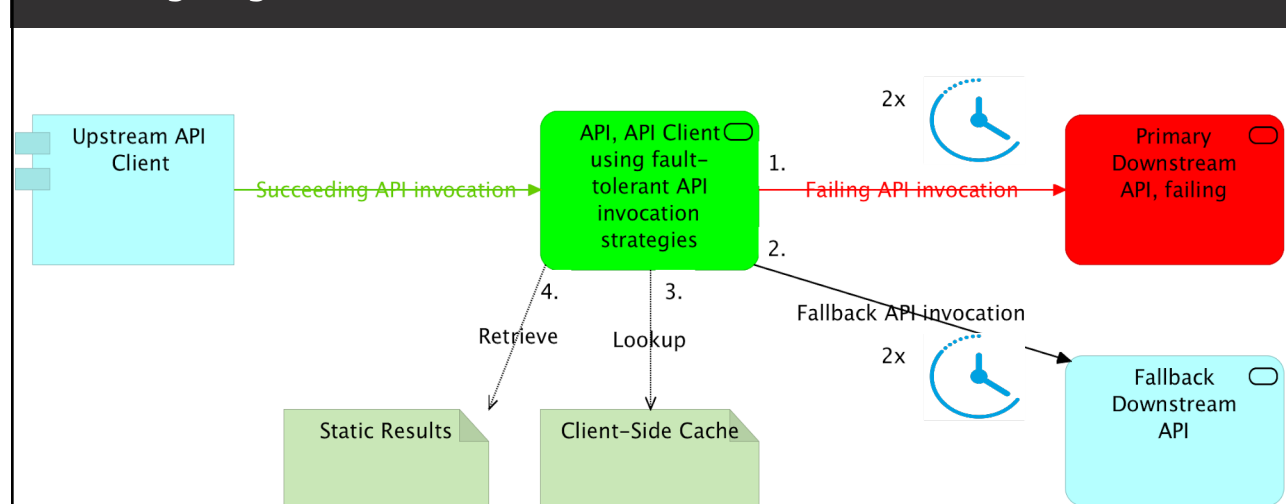
- **Prepared result** may be used instead of the result expected from an API invocation
- Best for APIs that return **reference data**-like results:
 - countries
 - states
 - currencies
 - products
- E.g., instead of result from "Policy Options Retrieval SAPI" a list of **common policy options** loaded from a configuration file:
 - **Not ideal** but better than not creating a quote at all
- For **Mule apps**:
 - Many options for results storage/retrieval, incl. **properties**

All contents © MuleSoft Inc.

43

43

Designing API clients for fault-tolerant API invocations



All contents © MuleSoft Inc.

44

44

Section 3

Understanding architecturally significant persistence choices for API implementations

45

Introducing CQRS as an API implementation strategy

- **CQRS**

- Command Query Responsibility Segregation
- Usage of **different models** for
 - **reading** from data (queries) and
 - **writing** to data (commands)

- Reads and writes can be optimized and scaled **independently**

- Independent **persistence** mechanisms for reads and writes

- **Commands**

- Formulated in the domain language of the **Bounded Context**
- Trigger **writes**

All content is typically queued and executed **asynchronously**

46

46

Introducing CQRS as an API implementation strategy

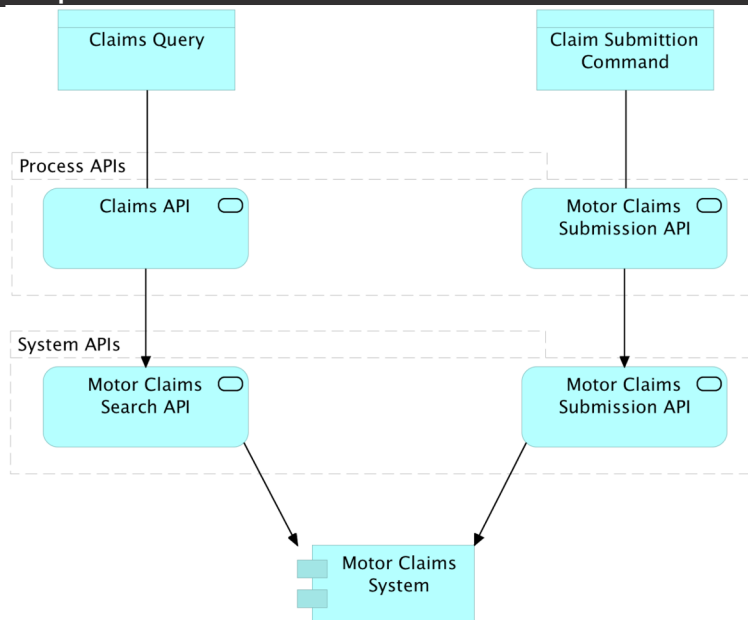
- **Queries**
 - Optimized for the **API clients'** needs (joins, aggregates)
 - Execute **synchronously**
 - May return slightly **out-of-date data**
- **Complicates** architecture and implementation of an API
- **Visible in API** itself
- Independent design choice **for each API implementation**
 - Architecturally **significant** because
 - Causes **eventual consistency** between read-sides and write-sides
 - Apparent in the **API specification**
 - May cause **API** to be **split** into one API for queries and one for commands

All contents © MuleSoft Inc.

47

47

Viewing Acme Insurance's application network from a CQRS perspective



All contents © MuleSoft Inc.

48

48

Introducing persistence with Event Sourcing



- **Event Sourcing**
 - Approach to data **persistence** that keeps persistent state as a series of **events** rather than as a snapshot of the current state
 - Often combined with CQRS
- Similar to **database transaction logs**, but in **application layer**
- **Events**
 - Expressed in the domain language of the **Bounded Context**
 - Typically arise directly from executing CQRS-style **commands**
 - **Propagate state changes** from CQRS write-side to CQRS read-side
 - Relies on **asynchronous messaging** systems like Anypoint MQ
- **Invisible in API specification: implementation detail**

All contents © MuleSoft Inc.

49

49

Summary



50

Summary



- Anypoint Platform-managed API implementations can target **Mule runtime or other runtimes**
- API implementations implemented as Mule apps can be **automatically discovered** by Anypoint Platform
- Anypoint Platform has over 120 **Connectors**
 - Indispensable for implementing **System APIs**
- **CloudHub** is an AWS-based iPaaS for the scalable, performant and highly-available deployment of Mule apps
- **Object Store** is a key-value persistence service available in all Mule runtime deployment scenarios

All contents © MuleSoft Inc.

51

51

Summary



- **In CloudHub**, Object Store allows time-limited persistence local to the AWS region of the Mule runtime
- API clients must guard against **API invocation failures**
 - In particular those which are API implementations
 - **Retry, Circuit Breaker, Fallbacks, ...**
- Some API implementations may benefit from **CQRS**
- **Separation of commands and queries** often arises naturally in API design, even in absence of true CQRS
- **Event Sourcing** is an implementation-level design decision of each API implementation

All contents © MuleSoft Inc.

52

52